

Wang_Miaohan_HW5

March 1, 2020

```
[1]: import random
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import BaggingClassifier, RandomForestClassifier,
↳ GradientBoostingClassifier
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import cross_val_score, GridSearchCV
from sklearn.inspection import plot_partial_dependence
```

Conceptual Exercise

When we have only two classes of outcome (supposedly outcome 1 and outcome 2), we will only have two probabilities \hat{p}_{m1} and \hat{p}_{m2} in our


```
[2]: # calculation of error rates in a simple two-class case
def c_error(p):
    return 1 - max(p, 1 - p)
def gini(p):
    return 2 * p * (1 - p)
def cross(p):
    return -(p * np.log(p) + (1 - p) * np.log(1 - p))

# generating plot
x_range = np.arange(0.01, 1, 0.01)
classification_error = list(map(c_error, x_range))
gini_index = list(map(gini, x_range))
cross_entropy = list(map(cross, x_range))

plt.figure(figsize=(9,7))
plt.plot(x_range, cross_entropy, label='cross entropy', color='tab:green')
plt.plot(x_range, gini_index, label='gini_index', color='tab:orange')
plt.plot(x_range, classification_error, label='classification error')
```

```
plt.xlabel('$\hat{p}_{m1}$', fontsize=15)
plt.ylabel('Error Rate', fontsize=15)
plt.title('Comparing Error Rate Methods in a Two-class Classification Setting',
         ↪ fontsize=20)
plt.legend();
```

Comparing Error Rate Methods in a Two-class Classification Setting



When growing a tree, we want to find the feature that can best split the data. Given that features can perform very similarly at splits, we need a sensitive measure of prediction error to differentiate feature performance. Among the three error rates, cross entropy is the most sensitive to changes in \hat{p}_{m1} , hence the best error rates to use in growing trees.

When pruning a tree, we want to cut redundant features to obtain the model with best predictive accuracy. Therefore, we would like to obtain the optimal tree depth with least mistakes in classification. Hence, in pruning, we incorporate the classification error, which tell us about model accuracy at each step in the most straightforward way.

Application Exercise

```
[3]: train_df = pd.read_csv('data/gss_train.csv')
test_df = pd.read_csv('data/gss_test.csv')

y_train = train_df.colrac
```

```

X_train = train_df.drop('colrac', axis=1)
y_test = test_df.colrac
X_test = test_df.drop('colrac', axis=1)

seed = 1234
error_rates = {}

```

Logistic Regression

```

[4]: log_mod = LogisticRegression(random_state=seed, max_iter=1000)
log_err = 1 - cross_val_score(log_mod, X_train, y_train, scoring='accuracy',
    ↪cv=10).mean()

log_mod.fit(X_train, y_train)
log_auc = roc_auc_score(y_train, log_mod.predict(X_train))

error_rates['Logistic Regression'] = [log_err, log_auc]

```

Naive Bayes

```

[5]: nb_mod = GaussianNB()
nb_err = 1 - cross_val_score(nb_mod, X_train, y_train, scoring='accuracy',
    ↪cv=10).mean()

nb_mod.fit(X_train, y_train)
nb_auc = roc_auc_score(y_train, nb_mod.predict(X_train))

error_rates['Naive Bayes'] = [nb_err, nb_auc]

```

Elastic Net Regression

```

[6]: # hyperparameter tuning
el_mod = LogisticRegressionCV(cv=10, penalty='elasticnet', scoring='accuracy',
    ↪max_iter=2000,
                                solver='saga', random_state=seed, l1_ratios = np.
    ↪arange(0, 1.1, 0.1))

el_mod.fit(X_train, y_train)
el_mod.l1_ratio_

```

```

[6]: array([0.2])

```

```
[7]: # fitting best model
el_mod_best = LogisticRegression(penalty='elasticnet', solver='saga',
    ↳l1_ratio=0.2, max_iter=1500)
el_err = 1 - cross_val_score(el_mod_best, X_train, y_train, scoring='accuracy',
    ↳cv=10).mean()
el_err

el_mod_best.fit(X_train, y_train)
el_auc = roc_auc_score(y_train, el_mod_best.predict(X_train))

error_rates['Elastic Net Regression'] = [el_err, el_auc]
```

Decision Tree (CART)

```
[8]: # hyperparameter tuning
tree_params = {'criterion':['gini', 'entropy'],
               'max_depth': np.arange(1, 56),
               'min_samples_leaf': np.arange(1, 11)}
tree_mod = GridSearchCV(DecisionTreeClassifier(), tree_params,
    ↳scoring='accuracy', cv=10)
tree_mod.fit(X_train, y_train)
tree_mod.best_estimator_
```

```
[8]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                           max_depth=4, max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=5, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, presort='deprecated',
                           random_state=None, splitter='best')
```

```
[9]: # fitting best model
tree_mod_best = DecisionTreeClassifier(criterion='gini', max_depth=4,
    ↳min_samples_leaf=5)
tree_err = 1 - cross_val_score(tree_mod_best, X_train, y_train,
    ↳scoring='accuracy', cv=10).mean()

tree_mod_best.fit(X_train, y_train)
tree_auc = roc_auc_score(y_train, tree_mod_best.predict(X_train))

error_rates['Decision Tree'] = [tree_err, tree_auc]
```

Bagging

```
[12]: # hyperparameter tuning
bag_params= {'n_estimators': np.arange(100, 900, 100)}
bag_mod = GridSearchCV(BaggingClassifier(max_samples=0.6, # suggested by https://
    ↪/www.youtube.com/watch?v=2Mg8QDOF1dQ
                                random_state=seed), bag_params,
    ↪scoring='accuracy', cv=10)
bag_mod.fit(X_train, y_train)
bag_mod.best_estimator_
```

```
[12]: BaggingClassifier(base_estimator=None, bootstrap=True, bootstrap_features=False,
                        max_features=1.0, max_samples=0.6, n_estimators=700,
                        n_jobs=None, oob_score=False, random_state=1234, verbose=0,
                        warm_start=False)
```

```
[13]: # fitting best model
bag_mod_best = BaggingClassifier(n_estimators=700, max_samples=0.6,
    ↪random_state=seed)
bag_err = 1 - cross_val_score(bag_mod_best, X_train, y_train,
    ↪scoring='accuracy', cv=10).mean()

bag_mod_best.fit(X_train, y_train)
bag_auc = roc_auc_score(y_train, bag_mod_best.predict(X_train))

error_rates['Bagging'] = [bag_err, bag_auc]
```

Random Forest

```
[19]: # hyperparameter tuning
rf_params= {'n_estimators': np.arange(100, 600, 100),
            'min_samples_leaf': np.arange(5, 11),
            'max_features': ['sqrt', 'log2']}
rf_mod = GridSearchCV(RandomForestClassifier(random_state=seed), rf_params,
    ↪scoring='accuracy', cv=10)
rf_mod.fit(X_train, y_train)
rf_mod.best_estimator_
```

```
[19]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                           criterion='gini', max_depth=None, max_features='sqrt',
                           max_leaf_nodes=None, max_samples=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=8, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=200,
                           n_jobs=None, oob_score=False, random_state=1234,
```

```
verbose=0, warm_start=False)
```

```
[20]: # fitting best model
rf_mod_best = RandomForestClassifier(n_estimators=200, min_samples_leaf=8,
    ↳max_features='sqrt', random_state=seed)
rf_err = 1 - cross_val_score(rf_mod_best, X_train, y_train, scoring='accuracy',
    ↳cv=10).mean()

rf_mod_best.fit(X_train, y_train)
rf_auc = roc_auc_score(y_train, rf_mod_best.predict(X_train))

error_rates['Random Forest'] = [rf_err, rf_auc]
```

Boosting

```
[22]: # hyperparameter tuning
bo_params = {'learning_rate': [0.001, 0.01],
            'n_estimators': [10, 100, 1000],
            'max_depth': [1, 2, 3]}
bo_mod = GridSearchCV(GradientBoostingClassifier(random_state=seed), bo_params,
    ↳cv=10, scoring='accuracy')
bo_mod.fit(X_train, y_train)
bo_mod.best_estimator_
```

```
[22]: GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse', init=None,
                                learning_rate=0.01, loss='deviance', max_depth=2,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=1000,
                                n_iter_no_change=None, presort='deprecated',
                                random_state=1234, subsample=1.0, tol=0.0001,
                                validation_fraction=0.1, verbose=0,
                                warm_start=False)
```

```
[23]: # fitting best model
bo_mod_best = GradientBoostingClassifier(learning_rate=0.01, n_estimators=1000,
    ↳max_depth=2, random_state=seed)
bo_err = 1 - cross_val_score(bo_mod_best, X_train, y_train, scoring='accuracy',
    ↳cv=10).mean()

bo_mod_best.fit(X_train, y_train)
bo_auc = roc_auc_score(y_train, bo_mod_best.predict(X_train))

error_rates['Boosting'] = [bo_err, bo_auc]
```

Model Evaluation

```
[24]: errors_df = pd.DataFrame(error_rates, index=['10-fold CV Error Rate', 'ROC/
      ↪AUC']).T
      errors_df.sort_values(by='10-fold CV Error Rate')
```

```
[24]:
```

	10-fold CV Error Rate	ROC/AUC
Random Forest	0.197844	0.870774
Boosting	0.202579	0.839971
Elastic Net Regression	0.203245	0.820179
Logistic Regression	0.203916	0.815254
Bagging	0.205304	0.992426
Decision Tree	0.219549	0.790820
Naive Bayes	0.265577	0.743125

Given that random forest has the lowest 10-fold cv error rate and the second highest ROC/AUC score, we will choose random forest as the best model for the gss data. Although random forest model does not yield the best ROC/AUC score, its ability to differentiate 0's as 0 and 1's as 1 is well enough among our selection of models. Bagging model performs with the greatest ROC/AUC score but it suffers from high variance, which gives it a less desirable error rate. Boosting also performs with low error rate, but it suffers from a lower ROC/AUC score, which means it makes more false positive and negatives in prediction. Also, boosting has a higher risk of overfitting, hence, we choose random forest as the best model to be implemented.

```
[30]: test_err = 1 - cross_val_score(rf_mod_best, X_test, y_test, scoring='accuracy',
      ↪cv=10).mean()
      test_auc = roc_auc_score(y_test, rf_mod_best.predict(X_test))

      print(f'When applied to the test set, the best random forest model yields
      ↪10-fold cv error rate of {test_err} \nand ROC/AUC score of {test_auc}')
```

When applied to the test set, the best random forest model yields 10-fold cv error rate of 0.20681632653061222 and ROC/AUC score of 0.7826961271102284


The random forest model falls in test error rate by only about 0.008, which is still lower than the training error of decision tree and naive bayes, hence maintains good predictive ability. The ROC/AUC score falls about 0.09 from train to test dataset. Given that the score is within 0.7-0.8, the random forest model maintains a fair power of avoiding false positives/negatives. Therefore, we can conclude that the random forest model does have good generalizability.

Partial Dependence Plots

```
[33]: features = ['tolerance', 'age']
      plot_partial_dependence(rf_mod_best, X_train, features,
                              n_jobs=3, grid_resolution=20)
      fig = plt.gcf()
```

```
fig.suptitle('Partial dependence of colrac on tolerance and age with Gradient_
↳Boosting Machine')
fig.subplots_adjust(wspace=0.4, hspace=0.3)
```

Partial dependence of colrac on tolerance and age with Gradient Boosting Machine



The PDPs shows that in our random forest model, there exists an negative non-linear relationship between political tolerance and the probability that the racist professor is allowed to teach. The higher value the tolerance variable takes, the less likely the professor is deemed “allowed to teach.” The slope of the PDP on tolerance is very steep, indicating that the tolerance variable has great influence on the prediction of colrac. The PDP of age, however, has a weak, negative, and seemingly linear relationship with the prediction. This shows that age has negative effect on colrac but not as strong as tolerance. The higher the professor’s age, the less likely the professor is allowed to teach.

(p.s. I have trouble finding the annotation of the ‘tolerance’ variable in GSS Data Explorer...therefore I can’t interpret the tolerance PDP concretely.)

```
[31]: rf_result = pd.DataFrame({'feature': X_train.columns, 'importance': rf_mod_best.
↳feature_importances_})
rf_result.sort_values(by='importance', ascending=False, inplace=True)
rf_result.head(10)
```

```
[31]:
```

	feature	importance
32	tolerance	0.241531
10	colmslm	0.145349
6	colath	0.101771
8	colmil	0.089509
7	colcom	0.042849
0	age	0.030211

12	egalit_scale	0.022110
18	income06	0.020313
23	pray	0.019338
26	science_quiz	0.016042

Looking at the feature importances in the random forest model, tolerance has the highest importance score, confirming its steep slope in PDP. The age variable has a much lower feature importance score, but there is still an effect, it's the seventh importance predictor in the random forest model.

[]: